

# Chunking Strategies for Retrieval-Augmented Generation

## Abstract

We study how chunk size and overlap affect retrieval precision and downstream answer quality across five document types. Smaller chunks improve precision but hurt recall on questions that span multiple chunks.

## Method

We index each corpus at chunk sizes of 200, 500, and 1000 tokens with 0% and 10% overlap, then evaluate retrieval against a held-out question set using a fixed embedding model.

## Results

500-token chunks with 10% overlap gave the best F1 across corpora. Overlap mattered more for narrative text than for structured documents.

## Discussion

Practitioners should tune chunk size per corpus rather than adopting a single default; the cost of re-indexing is low relative to the quality gain.